

Practical-7

Name : Yash Marthak
Roll No. : 19BCE122
Couse Name : Big Data Analytics

Aim:

Implement any one of the analytic algorithms using MapReduce by handling larger datasets in main memory.PCY/Multi-Hash/SON algorithm Regression K-means Clustering.

Here I have implemented the K-Means algorithm as given below.

K-means on MapReduce

k-means::Map

Input: Data points D , number of clusters k and centroids

1: for each data point $d \in D$ do

2: Assign d to the closest centroid

Output: centroids with associated data points

k-means::Reduce

Input: Centroids with associated data points

1: Compute the new centroids by calculating the average of data points in cluster

2: Write the global centroids to the disk

Output: New centroids

Code:

```
import java.io.*;
import java.util.*;
import java.util.StringTokenizer;

import org.apache.hadoop.conf.Configuration;
import org.apache.hadoop.fs.Path;
import org.apache.hadoop.fs.FSDataOutputStream;
import org.apache.hadoop.fs.FileSystem;
import org.apache.hadoop.io.IntWritable;
import org.apache.hadoop.io.Text;
```

```
import org.apache.hadoop.mapreduce.Job;
import org.apache.hadoop.mapreduce.Mapper;
import org.apache.hadoop.mapreduce.Reducer;
import org.apache.hadoop.mapreduce.lib.input.FileInputFormat;
import org.apache.hadoop.mapreduce.lib.output.FileOutputFormat;
```

```
public class Practical7_19bce277 {

    static List<Integer> Centroids;
    static List<Integer> newCentroids;
    static Path datapoints;
    static Path centroids;
    static FileSystem fs;

    public static class TokenizerMapper
        extends Mapper<Object, Text, Text, IntWritable>{

        private Text word = new Text();

        public void map(Object key, Text value, Context context
            ) throws IOException, InterruptedException {
            StringTokenizer itr = new StringTokenizer(value.toString());
            while (itr.hasMoreTokens()) {
                word.set(itr.nextToken());
                int datapoint = Integer.parseInt(word.toString());
                int min = 100000000;
                int center = 100000000;
                for(int i=0;i<Centroids.size();i++) {
                    int dist = Math.abs(datapoint-Centroids.get(i));
                    if(dist<min) {
                        min=dist;
                        center = Centroids.get(i);
                    }
                }
                context.write(new Text(""+center), new IntWritable(datapoint));
            }
        }
    }

    public static class IntSumReducer
        extends Reducer<Text,IntWritable,Text,IntWritable> {
        private IntWritable result = new IntWritable();

        public void reduce(Text key, Iterable<IntWritable> values,
            Context context
            ) throws IOException, InterruptedException {
            int sum = 0;
            int length = 0;
```

```

        for (IntWritable val : values) {
            sum += val.get();
            length++;
        }
        result.set(sum/length);
        newCentroids.add(Integer.parseInt(result.toString()));
    }
}

```

```

public static void main(String[] args) throws Exception {
    for(int i=0;i<3;i++) {
        Configuration conf = new Configuration();
        Centroids = new ArrayList<>();
        newCentroids = new ArrayList<>();
        datapoints=new Path("hdfs://19bce277_first/datapointsPractical7");//Location of file
in HDFS
        centroids=new Path("hdfs://19bce277_first/centroidsPractical7");//Location of file in
HDFS

```

```

        fs = FileSystem.get(conf);
        BufferedReader br=new BufferedReader(new InputStreamReader(fs.open(centroids)));
        String line;
        line=br.readLine();
        while (line != null){
            Centroids.add(Integer.parseInt(line));
            line=br.readLine();
        }


        Job job = Job.getInstance(conf, "word count");
        job.setJarByClass(Practical7_19bce277.class);
        job.setMapperClass(TokenizerMapper.class);
        job.setReducerClass(IntSumReducer.class);
        job.setOutputKeyClass(Text.class);
        job.setOutputValueClass(IntWritable.class);
        FileInputFormat.addInputPath(job, datapoints);
        FileOutputFormat.setOutputPath(job, new
Path("hdfs://19bce277_first/Practical7Demo"+i));
        if(job.waitForCompletion(true)) {
            FSDataOutputStream out = fs.create(centroids, true);
            BufferedWriter bw = new BufferedWriter(new OutputStreamWriter(out));
            for(Integer itr: newCentroids) {
                System.out.println(itr);
                bw.write(itr+"\n");
            }
            bw.close();
        }
    }
}
}

```

Now, I use two files:-

- 1) Stores the datapoints.
- 2) Stores the centroids.


Initially Centroids file looks like this:-

 centroidsPractical7.txt - Notepad

File Edit Format View Help

1
2
3

Datapoints:-

 datapointsPractical7.txt - Notepad

File Edit Format View Help

12
5
11
13
3
14
9
6
12
7
13
16
5
13
20
6
7
11
18
1

Running the MapReduce Job: -

```
2022-11-07 13:30:09,531 INFO impl.MetricsConfig: Loaded properties from hadoop-metrics2.properties
2022-11-07 13:30:09,591 INFO impl.MetricsSystemImpl: Scheduled Metric snapshot period at 10 second(s)
2022-11-07 13:30:09,591 INFO impl.MetricsSystemImpl: JobTracker metrics system started
2022-11-07 13:30:09,963 WARN mapreduce.JobResourceUploader: Hadoop command-line option parsing not p
the Tool interface and execute your application with ToolRunner to remedy this.
2022-11-07 13:30:10,047 INFO input.FileInputFormat: Total input files to process : 1
2022-11-07 13:30:10,157 INFO mapreduce.JobSubmitter: number of splits:1
2022-11-07 13:30:10,275 INFO mapreduce.JobSubmitter: Submitting tokens for job: job_local1580269013_
2022-11-07 13:30:10,275 INFO mapreduce.JobSubmitter: Executing with tokens: [1
```

In between I have printed the intermediate centroids:-

```
WRONG_REDUCE=0
File Input Format Counters
  Bytes Read=69
File Output Format Counters
  Bytes Written=0
1
10
```

Where 1 and 10 are the centroids after the 1st iteration.

```
Bytes Written=0
3
11
2022-11-07 13:30:13,036 WARN impl.MetricsSystemImpl: JobTracker metrics system already initialized!
2022-11-07 13:30:13,041 WARN mapreduce.JobResourceUploader: Hadoop command-line option parsing not performed. Implement
the Tool interface and execute your application with ToolRunner to remedy this.
2022-11-07 13:30:13,058 INFO input.FileInputFormat: Total input files to process : 1
2022-11-07 13:30:13,070 INFO mapreduce.JobSubmitter: number of splits:1
```

Where 3 and 11 are the centroids after the 2nd iteration.

```
File Output Format Counters
  Bytes Written=0
13
5
```

Where 13 and 5 are the centroids after the 3rd iteration.

The code is run just a single time for 3 iterations of MapReduce job.

Finally the centroids file looks as below:-

[Download](#)[Head the file \(first 32K\)](#)[Tail the file \(last 32K\)](#)

Block information —

Block 0 ▾

Block ID: 1073741912

Block Pool ID: BP-1112790740-172.23.224.1-1663568697260

Generation Stamp: 1088

Size: 5

Availability:

- N511-105.mshome.net

File contents

13

5